

CIS 552 DATABASE DESIGN

Optimizing Database Performance with Enhanced Hash Join Algorithms

Saheer Shaik
02080047

Likhil Naik Vislavath
02075503

Pruthvi Dhanavath
02082040

Introduction and Objectives:

The project, "Optimizing Database Performance with Enhanced Hash Join Algorithms," is driven by a compelling need to enhance database efficiency, a cornerstone in the realm of data management. At its core, the project is an in-depth exploration and refinement of the Hash Join algorithm, a technique integral to efficient database query processing. The significance of join operations in databases cannot be overstated, as they are pivotal in a wide array of applications, from intricate business intelligence analysis to crucial transaction processing systems.

Join operations, particularly equi-joins, form the backbone of relational database query processing, enabling the seamless merging and analysis of related data sets. These operations are fundamental in constructing meaningful relationships across different data tables, which is essential for data-driven decision-making and insights. However, the performance of these join operations, especially in databases handling voluminous or complex data sets, often poses challenges. These challenges can manifest as slower query response times or increased computational resource consumption, directly impacting the efficiency and scalability of database systems.

In this project, the focus is to tackle these performance challenges by optimizing the Hash Join algorithm, renowned for its efficiency in equi-join operations. The project aims to address specific issues that arise in scenarios involving large-scale or skewed data distributions, which are increasingly common in today's data-centric world. By enhancing the Hash Join algorithm and incorporating advanced hashing functions like XXHash and FarmHash, the project endeavours to improve the algorithm's performance, making it more robust and efficient in handling diverse data sets.

The project's objectives are twofold: firstly, to achieve a substantial improvement in the execution speed of the Hash Join algorithm, targeting a 20% increase in efficiency. Secondly, to determine the most effective hashing method under varying data conditions, providing a nuanced understanding of how different hashing functions impact join operation performance. This comprehensive approach, blending theoretical exploration with practical application, promises to significantly advance the field of database management, offering more efficient solutions for handling complex queries and large data sets.

In summary, "Optimizing Database Performance with Enhanced Hash Join Algorithms" is not just an academic exercise but a pragmatic endeavour aimed at bringing tangible improvements in database performance. It stands at the intersection of theoretical research and practical application, with the potential to yield significant benefits for a wide range of data-driven applications.

Database Design and Implementation:

The foundation of our project, "Optimizing Database Performance with Enhanced Hash Join Algorithms," lies in the meticulous design and implementation of our database, primarily structured around two central tables: orders and lineitem. These tables were constructed using data generated by TPC-H's dbgen tool, a benchmarking utility that produces realistic datasets commonly used for testing database performance.

Data Generation and Collection Process

We began the data generation process using TPC-H's dbgen tool. This tool is specifically designed to create test data for the TPC-H benchmark, simulating real-world database usage in a controlled environment. We selected this tool for its reliability and industry acceptance in generating standardized test data, ensuring our results would be relevant and comparable to real-world scenarios.

Steps to download TPC-H data:

1. Go to TPC-H official website → Downloads - https://www.tpc.org/tpc_documents_current_versions/current_specifications5.asp
2. Download [TPC-H Tools v3.0.1.zip](#)
3. Extract the file into local machine.
4. Open command prompt and go to dbgen directory from TPC-H tools folder.
5. Inside the directory, you will find a **makefile**. You might need to modify this file to suit your system's configuration (like the database and OS).
6. Compile the tool by running **make**. This should create an executable named **dbgen**.
7. Execute **./dbgen** to generate the data.
8. You can specify various options like **-s [scale factor]** to determine the volume of data generated. For example, **-s 1** generates a 1GB dataset.
9. The execution of **dbgen** will result in several **.tbl** files, each corresponding to a TPC-H table (e.g., **customer.tbl**, **orders.tbl**, **lineitem.tbl**, etc.).
10. If needed, you can convert the **.tbl** files into other formats like CSV for easier importing into SQL databases.

After installing and configuring dbgen, we generated the datasets for our primary tables: orders and lineitem. The process involved specifying parameters like scale factor and data format, which allowed us to tailor the data volume and structure to our project needs. The generated data was then exported as CSV files, which we later imported into our SQL database environment for further manipulation and analysis.

The orders Table

The orders table is a central component of our database schema. Its structure is as follows:

- O_ORDERKEY: INTEGER, PRIMARY KEY. This is the unique identifier for each order.
- O_CUSTKEY: INTEGER, FOREIGN KEY. Links to the customer table (not part of our immediate scope).

- O_ORDERSTATUS: CHAR(1). Represents the status of the order (e.g., 'O' for open, 'F' for complete).
- O_TOTALPRICE: DECIMAL(15,2). The total price of the order.
- O_ORDERDATE: DATE. The date when the order was placed.
- O_ORDERPRIORITY: CHAR(15). Indicates the priority of the order.
- O_CLERK: CHAR(15). The clerk handling the order.
- O_SHIPPRIORITY: INTEGER. A numerical value indicating the priority of shipment.
- O_COMMENT: VARCHAR(79). Additional comments about the order.

The lineitem Table

The lineitem table details individual items within each order:

- L_ORDERKEY: INTEGER, PRIMARY KEY (Composite with L_LINENUMBER), FOREIGN KEY (references O_ORDERKEY in orders). Identifies the order to which the line item belongs.
- L_PARTKEY: INTEGER. Identifies the part number.
- L_SUPPKEY: INTEGER. Identifies the supplier key.
- L_LINENUMBER: INTEGER, PRIMARY KEY (Composite with L_ORDERKEY). Line number in the order.
- L_QUANTITY: DECIMAL(15,2). Quantity of the item ordered.
- L_EXTENDEDPRICE: DECIMAL(15,2). Extended price of the line item.
- L_DISCOUNT: DECIMAL(15,2). Discount on the line item.
- L_TAX: DECIMAL(15,2). Tax on the line item.
- L_RETURNFLAG: CHAR(1). Return status of the line item.
- L_LINestatus: CHAR(1). Status of the line item.
- L_SHIPDATE: DATE. Date when the item was shipped.
- L_COMMITDATE: DATE. Date when the order was committed.
- L_RECEIPTDATE: DATE. Date when the item was received.
- L_SHIPINSTRUCT: CHAR(25). Shipping instructions.
- L_SHIPMODE: CHAR(10). Mode of shipping.
- L_COMMENT: VARCHAR(44). Additional comments.

Rationale and Relevance

The rationale behind selecting these two tables for our project is rooted in their intrinsic relationship and their relevance in typical business scenarios. The orders table represents high-level order information, while the lineitem table provides granular details about each item within those orders. This relational structure is emblematic of typical database schemas found in business environments, where data is segmented yet interconnected.

By focusing on these two tables, our project aims to address real-world database performance issues. The relationship between orders and lineitem necessitates efficient join operations for queries that span across these tables – a common scenario in database analytics and transaction processing. Our optimization strategies, therefore, have direct applicability in improving the performance of such operations, which are crucial in database environments handling large volumes of transactional data.

SQL Queries and Optimization Techniques:

In our project, we crafted SQL queries and employed strategic optimization techniques to improve database performance, particularly focusing on the orders and lineitem tables.

Loading Data

Initially, we encountered a challenge in loading data into our SQL database. Directly loading data from local files was restricted due to MySQL's security settings, which often include the `--secure-file-priv` option. This option limits the directories from which MySQL can load files as a security measure. To bypass this restriction, we utilized MySQL Workbench's Import Wizard. This tool provides a graphical interface to import data, seamlessly handling file path restrictions and ensuring a smooth data import process.

Detailed Walkthrough of SQL Queries

1. Database and Table Creation:

- The **CREATE DATABASE** and **CREATE TABLE** statements were used to establish the foundational structure of our database. These statements were crucial for defining the schema and ensuring the data integrity of the **orders** and **lineitem** tables.

```
1 CREATE DATABASE IF NOT EXISTS tpch;
2 USE tpch;
3
4 CREATE TABLE orders (
5     O_ORDERKEY    INTEGER NOT NULL,
6     O_CUSTKEY     INTEGER NOT NULL,
7     O_ORDERSTATUS CHAR(1) NOT NULL,
8     O_TOTALPRICE  DECIMAL(15,2) NOT NULL,
9     O_ORDERDATE   DATE NOT NULL,
10    O_ORDERPRIORITY CHAR(15) NOT NULL,
11    O_CLERK        CHAR(15) NOT NULL,
12    O_SHIPPRIORITY INTEGER NOT NULL,
13    O_COMMENT      VARCHAR(79) NOT NULL
14 );
15
16 CREATE TABLE lineitem (
17     L_ORDERKEY    INTEGER NOT NULL,
18     L_PARTKEY     INTEGER NOT NULL,
19     L_SUPPKEY     INTEGER NOT NULL,
20     L_LINENUMBER  INTEGER NOT NULL,
21     L_QUANTITY    DECIMAL(15,2) NOT NULL,
22     L_EXTENDEDPRICE DECIMAL(15,2) NOT NULL,
23     L_DISCOUNT   DECIMAL(15,2) NOT NULL,
24     L_TAX         DECIMAL(15,2) NOT NULL,
25     L_RETURNFLAG  CHAR(1) NOT NULL,
26     L_LINESTATUS  CHAR(1) NOT NULL,
27     L_SHIPDATE    DATE NOT NULL,
28     L_COMMITDATE  DATE NOT NULL,
29     L_RECEIPTDATE DATE NOT NULL,
30     L_SHIPINSTRUCT CHAR(25) NOT NULL,
31     L_SHIPMODE     CHAR(10) NOT NULL,
32     L_COMMENT     VARCHAR(44) NOT NULL
33 );
```

2. Primary and Foreign Key Constraints:

- ALTER TABLE** statements were used to add primary and foreign keys, crucial for maintaining relational integrity and optimizing join operations.

```
35 -- Add primary keys
36 ALTER TABLE orders ADD PRIMARY KEY (O_ORDERKEY);
37 ALTER TABLE lineitem ADD PRIMARY KEY (L_ORDERKEY, L_LINENUMBER);
38
39 -- Add foreign keys
40 ALTER TABLE lineitem ADD FOREIGN KEY (L_ORDERKEY) REFERENCES orders (O_ORDERKEY);
```

Indexing Strategies

Creation of Indexes:

We created indexes on frequently joined and queried columns such as O_ORDERKEY and L_ORDERKEY. This strategy significantly improved query performance by reducing the time complexity of search operations within these tables.

```
43 CREATE INDEX idx_orders_orderkey ON orders (O_ORDERKEY);
44 CREATE INDEX idx_lineitem_orderkey ON lineitem (L_ORDERKEY);
45 SELECT *
46 FROM orders
47 JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY;
```

Impact of Indexing on Performance:

The addition of indexes proved to be a game-changer, especially noticeable in complex join operations where the search time was substantially reduced.

Duration/Fetch time before and after optimization using indexing:

```
31 23:43:15 SELECT * FROM orders JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_OR... 1000 row(s) returned 0.0014 sec / 0.0047 s...
```

```
34 00:29:09 SELECT * FROM orders JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_OR... 1000 row(s) returned 0.013 sec / 0.0038 sec
```

Understanding the Metrics

Duration: This is the time taken to execute the query, including parsing, optimizing, and actually running it.

Fetch Time: This refers to the time taken to retrieve the results from the database and make them available for use.

Analysis of Results

1. Before Optimization:

- Duration: 0.0014 seconds
- Fetch Time: 0.0047 seconds
- These initial metrics suggest that while the query executed fairly quickly, fetching the results took comparatively longer.

2. After Indexing:

- Duration: 0.013 seconds
- Fetch Time: 0.0038 seconds
- Post-indexing, we see a slight increase in the duration but a decrease in the fetch time.

Interpretation and Significance

1. **Increased Duration:** The increase in execution duration might initially seem counterintuitive. However, this can be attributed to the overhead of using indexes. When a query utilizes an index, the database engine must first read the index and then use it to find the actual data. This additional step can sometimes add a small amount of overhead, hence the slightly longer duration.
2. **Reduced Fetch Time:** The decrease in fetch time is more in line with our expectations from indexing. With the relevant data more efficiently organized and quickly accessible through the index, the time to fetch and return this data is reduced. This is particularly significant in scenarios where the same data is queried frequently.

Why Every Millisecond Matters

- **Scalability:** In a large-scale database handling millions of queries, even minor improvements can accumulate to substantial time savings.
- **Resource Utilization:** Efficient fetch times mean reduced workload on the database server, leading to better resource utilization and the potential for handling more concurrent requests.
- **User Experience:** In applications where database response time is critical (like in web applications), even marginal decreases in response times can significantly enhance the user experience.

Optimization Techniques in Queries

1. **Optimizing JOIN Operations:**
 - In queries involving JOINS between **orders** and **lineitem**, we focused on optimizing the retrieval process by indexing the join keys. This reduced the time taken to merge data from these tables.
2. **Selective Data Retrieval:**
 - We optimized SELECT queries to fetch only necessary columns, reducing the processing load and memory usage. This approach was particularly effective in queries where large datasets were involved.
3. **Use of WHERE Clauses:**
 - The application of targeted WHERE clauses enabled us to filter out irrelevant data early in the query process, leading to faster execution times.

```
49      -- Optimize the SELECT clause to fetch only necessary columns
50      SELECT orders.O_ORDERKEY, lineitem.L_ORDERKEY, orders.O_TOTALPRICE, lineitem.L_QUANTITY, lineitem.L_SHIPDATE
51      FROM orders
52      JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY
53      WHERE lineitem.L_SHIPDATE >= '1995-04-25' AND lineitem.L_SHIPDATE < '1996-02-01'
54      AND orders.O_TOTALPRICE > 50000;
```

Performance Benchmarking with TPC-H

In our project, we conducted a comprehensive performance benchmarking using the TPC-H (Transaction Processing Performance Council's Benchmark H) standard. This benchmark is specifically designed for evaluating the performance of various database operations, including complex queries and data manipulation.

Procedure for TPC-H Benchmark Testing

1. Setup:

- We used the **dbgen** tool from TPC-H to generate a realistic dataset. This dataset mimics real-world business data, making our benchmarking relevant and practical.
- The generated data was imported into our SQL database, creating the **orders** and **lineitem** tables as primary subjects of our testing.

2. Execution:

- TPC-H provides a suite of queries (known as Query 1 to Query 22), each designed to test different aspects of database performance.
- We selected four of these queries that best matched our project's focus on join operations and optimizations.

3. Data Handling:

- The data generated by **dbgen** was carefully handled to maintain its integrity and representativeness, ensuring that our benchmarking results would be accurate and meaningful.

Detailed Analysis of Benchmark Query Results

We focused on four specific TPC-H benchmark queries, optimizing each one. Below is a detailed analysis of these queries, both before and after our optimization efforts:

1. Query 1: Join Operation Optimization

- **Original Query:** Involves a left join between the **lineitem** and **orders** tables. The query was designed to test the database's capability to handle complex join operations.


```

58 SELECT l_count, COUNT(*) AS lineitemcount
59 FROM (
60     SELECT
61         l.L_ORDERKEY,
62         COUNT(o.O_ORDERKEY)
63     FROM
64         lineitem AS l
65     LEFT OUTER JOIN orders AS o ON l.L_ORDERKEY = o.O_ORDERKEY
66     AND o.O_COMMENT NOT LIKE '%[WORD1]%' '[WORD2]%'
67     GROUP BY
68         l.L_ORDERKEY
69 ) AS l_orders (L_ORDERKEY, l_count)
70 GROUP BY
71     l_count
72 ORDER BY
73     lineitemcount DESC,
74     l_count DESC;

```

- **Optimization:** We created indexes on the join keys (**L_ORDERKEY** in **lineitem** and **O_ORDERKEY** in **orders**) to speed up the join operation. Additionally, we refined the query to filter results based on specific comment patterns.

```

82 -- Optimized Query
83 SELECT l_count, COUNT(*) AS lineitemcount
84 FROM (
85     SELECT
86         l.L_ORDERKEY,
87         COUNT(o.O_ORDERKEY) -- Counting orders for each line item
88     FROM
89         lineitem AS l
90     LEFT OUTER JOIN orders AS o ON l.L_ORDERKEY = o.O_ORDERKEY
91     AND o.O_COMMENT NOT LIKE '%ms integrate fluffily. ironic asymptotes after the regular excuses nag around husly final asymptotes %' -- Filtering based on comment
92     GROUP BY
93         l.L_ORDERKEY -- Grouping by line item order key
94 ) AS l_orders (L_ORDERKEY, l_count)
95 GROUP BY
96     l_count -- Grouping results by the count of orders
97 ORDER BY
98     lineitemcount DESC, l_count DESC; -- Ordering the final result set

```

- **Result:** Post-optimization, the query showed a noticeable improvement in execution time, demonstrating the effectiveness of index-based optimizations in join operations.

35	01:01:15	SELECT l_count, COUNT(*) AS lineitemcount FROM (SELECT l.L_ORDE...	7 row(s) returned	0.562 sec / 0.00014...
56	01:41:26	SELECT l_count, COUNT(*) AS lineitemcount FROM (SELECT l.L_ORDE...	7 row(s) returned	0.608 sec / 0.000012...

2. Query 2: Aggregation and Grouping

- **Original Query:** Focused on performing aggregations and grouping on the joined results of **orders** and **lineitem**.

```

100 -- Query2:
101 SELECT
102     l_shipmode,
103     SUM(CASE WHEN o_orderpriority IN ('1-URGENT', '2-HIGH') THEN 1 ELSE 0 END) AS high_line_count,
104     SUM(CASE WHEN o_orderpriority NOT IN ('1-URGENT', '2-HIGH') THEN 1 ELSE 0 END) AS low_line_count
105 FROM
106     orders
107 JOIN
108     lineitem ON o_orderkey = l_orderkey
109 WHERE
110     l_shipmode IN ('TRUCK', 'RAIL')
111     AND l_commitdate < l_receiptdate
112     AND l_shipdate < l_commitdate
113     AND l_receiptdate >= '1996-02-03'
114     AND l_receiptdate < DATE_ADD('1996-02-03', INTERVAL 1 YEAR)
115 GROUP BY
116     l_shipmode
117 ORDER BY
118     l_shipmode;

```

- **Optimization:** We optimized this query by adding indexes on columns used in the **WHERE** clause, like **L_SHIPMODE** and **L_RECEIPTDATE**, to speed up data filtering.

```

120 -- Refined Query2:
121
122 -- Creating indexes on join keys for efficient join operation
123 CREATE INDEX idx_orders_orderkey ON orders (O_ORDERKEY);
124 CREATE INDEX idx_lineitem_orderkey ON lineitem (L_ORDERKEY);
125
126 -- Creating indexes on columns used in WHERE clause for faster filtering
127 CREATE INDEX idx_lineitem_shipmode ON lineitem (L_SHIPMODE);
128 CREATE INDEX idx_lineitem_receiptdate ON lineitem (L_RECEIPTDATE);
129
130 -- Analyzing the execution plan of the query
131 EXPLAIN
132 SELECT
133     l.l_shipmode,
134     -- Using SUM with CASE to count based on order priority conditions
135     SUM(CASE WHEN o.o_orderpriority IN ('1-URGENT', '2-HIGH') THEN 1 ELSE 0 END) AS high_line_count,
136     SUM(CASE WHEN o.o_orderpriority NOT IN ('1-URGENT', '2-HIGH') THEN 1 ELSE 0 END) AS low_line_count
137 FROM
138     orders AS o
139 JOIN
140     lineitem AS l ON o.o_orderkey = l.l_orderkey
141 WHERE
142     -- Filtering conditions to narrow down the result set
143     l.l_shipmode IN ('TRUCK', 'RAIL')
144     AND l.l_commitdate < l.l_receiptdate
145     AND l.l_shipdate < l.l_commitdate
146     AND l.l_receiptdate >= '1996-02-03'
147     AND l.l_receiptdate < DATE_ADD('1996-02-03', INTERVAL 1 YEAR)
148 GROUP BY
149     -- Grouping results by ship mode
150     l.l_shipmode
151 ORDER BY
152     -- Ordering the final result set by ship mode
153     l.l_shipmode;

```

- **Result:** The query execution became faster, highlighting the importance of indexing on aggregation and grouping operations.

✓ 40	01:16:22	SELECT l_shipmode, SUM(CASE WHEN o_orderpriority IN ('1-URGENT', '2...	2 row(s) returned	0.453 sec / 0.000016...
✓ 52	01:29:48	SELECT l.l_shipmode, -- Using SUM with CASE to count based on order...	2 row(s) returned	0.568 sec / 0.000012...

3. Query 3: Complex View Creation

- **Original Query:** Involved creating a view to analyze order characteristics based on the count of line items.

```
155      -- Query3:
156      CREATE VIEW orders_analysis AS
157      SELECT
158          o.O_ORDERKEY,
159          COUNT(l.L_ORDERKEY) AS lineitemcount
160      FROM
161          orders AS o
162      LEFT OUTER JOIN lineitem AS l ON o.O_ORDERKEY = l.L_ORDERKEY
163      AND l.L_COMMENT NOT LIKE '%sleep quickly. req%lithely regular deposits. fluffily %'
164      GROUP BY
165          o.O_ORDERKEY;
166
167      SELECT
168          lineitemcount,
169          COUNT(*) AS ordercount
170      FROM
171          orders_analysis
172      GROUP BY
173          lineitemcount
174      ORDER BY
175          ordercount DESC,
176          lineitemcount DESC;
177
178      DROP VIEW orders_analysis;
```

- **Optimization:** We ensured that indexes were in place for all key columns used in the view creation and query execution.

```
180      -- Refined Query3:
181      -- Ensure indexing is done for efficient join operations
182      -- CREATE INDEX IF NOT EXISTS idx_orders_orderkey ON orders (O_ORDERKEY);(Index created already)
183      -- CREATE INDEX IF NOT EXISTS idx_lineitem_orderkey ON lineitem (L_ORDERKEY);(Index created already)
184
185      -- Creating a view to analyze order characteristics
186      CREATE OR REPLACE VIEW orders_analysis AS
187      SELECT
188          o.O_ORDERKEY,
189          COUNT(l.L_ORDERKEY) AS lineitemcount -- Counting line items per order
190      FROM
191          orders AS o
192      LEFT OUTER JOIN lineitem AS l ON o.O_ORDERKEY = l.L_ORDERKEY
193      AND l.L_COMMENT NOT LIKE '%sleep quickly. req%lithely regular deposits. fluffily %' -- Filtering based on line item comment
194      GROUP BY
195          o.O_ORDERKEY;
196
197      -- Query to get distribution of orders based on line item count
198      SELECT
199          lineitemcount,
200          COUNT(*) AS ordercount -- Counting orders for each line item count
201      FROM
202          orders_analysis
203      GROUP BY
204          lineitemcount
205      ORDER BY
206          ordercount DESC,
207          lineitemcount DESC;
208
209      -- Dropping the view after analysis
210      DROP VIEW IF EXISTS orders_analysis;
```

- **Result:** This optimization led to a faster view creation and query execution time, reinforcing the role of indexes in complex operations.

✓ 65	03:52:07	SELECT lineitemcount, COUNT(*) AS ordercount FROM orders_analysis...	7 row(s) returned	0.737 sec / 0.000011...
✓ 71	04:36:33	SELECT lineitemcount, COUNT(*) AS ordercount -- Counting orders for ea...	7 row(s) returned	0.707 sec / 0.000016...

Our optimization may have led to a more complex data retrieval process. Even though the query executed faster, the data prepared for fetching might be more complex or larger in size, leading to a longer fetch time.

The key objective of our optimization was to reduce the total time from query initiation to result delivery. The slight increase in fetch time is offset by the more significant reduction in execution duration, leading to an overall performance gain.

The optimizations made the query more efficient, as evidenced by the reduced duration. This efficiency is vital in high-load scenarios, where faster query execution can lead to better throughput and resource utilization.

4. Query 4: Advanced Filtering and Calculation

- **Original Query:** Required advanced filtering and calculation of revenue based on order status.

```

212      -- Query4:
213      SELECT
214      100.00 * SUM(CASE
215          WHEN o.o_orderstatus = 'O' THEN l.l_extendedprice * (1 - l.l_discount)
216          ELSE 0
217      END) /
218      SUM(l.l_extendedprice * (1 - l.l_discount)) AS revenue_O,
219      100.00 * SUM(CASE
220          WHEN o.o_orderstatus = 'F' THEN l.l_extendedprice * (1 - l.l_discount)
221          ELSE 0
222      END) /
223      SUM(l.l_extendedprice * (1 - l.l_discount)) AS revenue_F
224      FROM
225          lineitem AS l
226      JOIN
227          orders AS o ON l.l_orderkey = o.o_orderkey
228      WHERE
229          l.l_shipdate >= '1995-04-25' -- Replace with your start date
230          AND l.l_shipdate < DATE_ADD('1995-04-25', INTERVAL 1 MONTH) -- Replace with your start date

```

- **Optimization:** We simplified the query by focusing on one revenue calculation at a time and ensuring optimal use of indexes.

```

234 SELECT
235     100.00 * SUM(CASE
236         WHEN o.o_ORDERSTATUS = '0' THEN l.l_EXTENDEDPRICE * (1 - l.l_DISCOUNT)
237         ELSE 0
238     END) /
239     SUM(l.l_EXTENDEDPRICE * (1 - l.l_DISCOUNT)) AS revenue_percentage_0
240 FROM
241     lineitem AS l
242 JOIN
243     orders AS o ON l.l_ORDERKEY = o.o_ORDERKEY
244 WHERE
245     l.l_SHIPDATE >= '1995-04-25' -- Replace with your actual start date
246     AND l.l_SHIPDATE < DATE_ADD('1995-04-25', INTERVAL 1 MONTH) -- Replace with your actual start date
247

```

- **Result:** The simplified query exhibited a more efficient execution, highlighting the benefits of query simplification in complex calculations.

✓ 84	05:20:31	SELECT	100.00 * SUM(CASE	WHEN o.o_ORDERSTATUS = '0' T...	1 row(s) returned	0.211 sec / 0.000009...
✓ 86	15:29:06	SELECT	100.00 * SUM(CASE	WHEN o.o_ORDERSTATUS = '0' T...	1 row(s) returned	0.263 sec / 0.000017...

The optimization likely involved restructuring the query or modifying indexes, which, while possibly adding slight overhead to query processing (hence the increased duration), made the retrieval of data much more efficient. This efficiency is reflected in the significantly reduced fetch time.

Interpretation of Performance Variations and Insights Gained

The optimization of each query under the TPC-H benchmark yielded valuable insights:

Indexing is Key: The addition of indexes on join keys and frequently filtered columns significantly improved query performance.

Simplification Works: Simplifying complex queries and breaking them down into smaller parts can lead to better performance.

Tailored Optimization: Each query required a unique optimization approach, reflecting the diverse nature of real-world database usage.

Advanced Hashing Functions: XXHash and FarmHash

In our project, we ventured beyond traditional SQL-based optimizations to explore the integration of advanced hashing functions - specifically XXHash and FarmHash. These functions are renowned for their speed and efficiency in computing hash values, making them potentially valuable for database operations, especially hash joins. Our exploration was conducted using Python due to certain limitations and considerations in SQL environments.

1. Limitations in SQL:

- Standard SQL implementations typically do not offer direct support for advanced hashing algorithms like XXHash and FarmHash. While some database systems allow custom functions, they often come with restrictions and may not support external libraries efficiently.
- SQL's primary focus is on data manipulation and retrieval, not on implementing complex hashing algorithms, which are more computation-focused.

2. Flexibility and Control in Python:

- Python, with its extensive library ecosystem and its ability to interface with various systems, provided us with the necessary tools and flexibility to implement and test these advanced hashing functions.
- Python's ability to handle complex logic and its efficient data processing capabilities made it an ideal choice for this part of the project.

Exploration of XXHash and FarmHash

1. implementing Hash Joins in Python:

- We used Python to implement hash join algorithms that utilize XXHash and FarmHash for hashing. This involved creating Python functions that mimicked the behavior of a hash join as it would occur in a database system.
- Our implementation read data from CSV files generated by the TPC-H **dbgen** tool, processed it in Python, and then applied the hash join using these advanced hashing functions.

2. Why XXHash and FarmHash:

- XXHash is known for its extremely fast execution, making it a good candidate for operations where speed is crucial.
- FarmHash, developed by Google, is reputed for its high efficiency and low collision rates in hashing, which is vital for maintaining data integrity in join operations.

Performance Metrics:

- We focused on key performance metrics such as execution time, memory usage, and the rate of hash collisions.
- These metrics provided a quantitative basis to compare the effectiveness of XXHash and FarmHash in the context of hash joins.
- While XXHash excelled in speed, FarmHash demonstrated a slightly better balance between speed and collision handling.
- For instance, in scenarios where speed is the top priority, XXHash might be preferable. In contrast, for applications where data integrity and collision avoidance are critical, FarmHash could be more suitable.

Optimized XXhash

```
port csv
port time
port xxhash

f read_csv_data_in_chunks(file_path, chunk_size=10000):
    with open(file_path, 'r', newline='', encoding='utf-8') as csvfile:
        reader = csv.reader(csvfile)
        chunk = []
        for i, row in enumerate(reader):
            if i % chunk_size == 0 and i > 0:
                yield chunk
                chunk = []
            chunk.append(row)
        yield chunk

f xxhash_key(value):
    return xxhash.xxh64(value).intdigest()

f hash_join(orders_chunk, lineitem, hash_function):
    hash_table = {}
    for order in orders_chunk:
        hash_key = hash_function(order[0])
        hash_table[hash_key] = order

    joined_data = []
    for item in lineitem:
        hash_key = hash_function(item[0])
        if hash_key in hash_table:
            joined_data.append(hash_table[hash_key] + item)

    return joined_data

ders_path = '/Users/shaiksaheer/Documents/College_Docs/DataBaseDesign/V3.0.1/dbgen/Data
neitem_path = '/Users/shaiksaheer/Documents/College_Docs/DataBaseDesign/V3.0.1/dbgen/Da

art_read_time = time.time()
neitem_data = list(read_csv_data_in_chunks(lineitem_path, chunk_size=10000))[0]
ad_time = time.time() - start_read_time

art_join_time = time.time()
ined_data = []
r orders_chunk in read_csv_data_in_chunks(orders_path, chunk_size=10000):
    joined_data.extend(hash_join(orders_chunk, lineitem_data, xxhash_key))
in_time = time.time() - start_join_time

tput_path = '/Users/shaiksaheer/Documents/College_Docs/DataBaseDesign/V3.0.1/dbgen/Data
th open(output_path, 'w', newline='', encoding='utf-8') as csvfile:
    csvwriter = csv.writer(csvfile)
    for row in joined_data:
        csvwriter.writerow(row)

int(f'Read time: {read_time:.4f} seconds')
int(f'Join time: {join_time:.4f} seconds')
int(f'Total time: {read_time + join_time:.4f} seconds')
int(f'Joined data has been saved to {output_path}')

Read time: 2.7327 seconds
Join time: 0.7529 seconds
Total time: 3.4857 seconds
```

Optimized FarmHash

```
import csv
import time
import farmhash

def read_csv_data_in_chunks(file_path, chunk_size=10000):
    with open(file_path, 'r', newline='', encoding='utf-8') as csvfile:
        reader = csv.reader(csvfile)
        chunk = []
        for i, row in enumerate(reader):
            if i % chunk_size == 0 and i > 0:
                yield chunk
                chunk = []
            chunk.append(row)
        yield chunk

def farmhash_key(value):
    return farmhash.hash64(value)

def hash_join(orders_chunk, lineitem, hash_function):
    hash_table = {}
    for order in orders_chunk:
        hash_key = hash_function(order[0])
        hash_table[hash_key] = order

    joined_data = []
    for item in lineitem:
        hash_key = hash_function(item[0])
        if hash_key in hash_table:
            joined_data.append(hash_table[hash_key] + item)

    return joined_data

orders_path = '/Users/shaiksaheer/Documents/College_Docs/DataBaseDesign/V3.0.1/dbgen/Da
lineitem_path = '/Users/shaiksaheer/Documents/College_Docs/DataBaseDesign/V3.0.1/dbgen/Da

start_read_time = time.time()
lineitem_data = list(read_csv_data_in_chunks(lineitem_path, chunk_size=10000))[0] # Re

# Perform the hash join in chunks
start_join_time = time.time()
joined_data = []
for orders_chunk in read_csv_data_in_chunks(orders_path, chunk_size=10000):
    joined_data.extend(hash_join(orders_chunk, lineitem_data, farmhash_key))
join_time = time.time() - start_join_time

output_path = '/Users/shaiksaheer/Documents/College_Docs/DataBaseDesign/V3.0.1/dbgen/Da
with open(output_path, 'w', newline='', encoding='utf-8') as csvfile:
    csvwriter = csv.writer(csvfile)
    for row in joined_data:
        csvwriter.writerow(row)

read_time = time.time() - start_read_time
print(f'Read time: {read_time:.4f} seconds')
print(f'Join time: {join_time:.4f} seconds')
print(f'Total time: {read_time + join_time:.4f} seconds')
print(f'Joined data has been saved to {output_path}')

Read time: 5.4128 seconds
Join time: 0.6477 seconds
Total time: 6.0605 seconds
```

Comparing Python and SQL Times

- Performance: SQL's fetch times are generally faster for database operations due to the specialized nature of database management systems. Python's total time (read plus join) is more indicative of a data processing pipeline that includes data loading and processing steps.
- Use Case Suitability: Python's approach is more suited for scenarios where data is not already in a database, or for complex operations that go beyond standard SQL capabilities. SQL excels in pure database operations and is generally more efficient for straightforward data retrieval and joins.
- Flexibility vs. Efficiency: Python offers more flexibility in data processing and can handle a variety of data sources and formats. SQL provides a more efficient but less flexible environment, tailored specifically for database operations.
- SQL databases are specifically designed to handle join operations efficiently. They use sophisticated algorithms and can leverage database indexes.
- Python's join time, while optimized within the scope of Python's capabilities, may not match the efficiency of a database system optimized for such operations.

In summary, while Python provides a versatile environment for data manipulation and custom operations, SQL fetch times are typically shorter for database-specific tasks due to the inherent optimizations in database systems. The choice between Python and SQL should be guided by the specific requirements of the task at hand.

Addressing Key Project Questions:

Question 1: How Does the Introduction of Indexing Affect the Performance of Join Operations Between Orders and Lineitem Tables?

This question aims to assess the impact of indexing on the performance of join operations, a key aspect of your project.

```
-- The query before adding indexes
SELECT *
FROM orders
JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY;
```

```
-- The query after adding indexes
-- (Assuming indexes have been created on O_ORDERKEY and L_ORDERKEY)
SELECT *
FROM orders
JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY;
```

TRC Representation:

$$\{ t \mid \exists o \in \text{orders}, \exists l \in \text{lineitem} (o.O_ORDERKEY = l.L_ORDERKEY \wedge t = o \bowtie l) \}$$

Query Plan:

1. **Scan:** Scan both **orders** and **lineitem** tables.
2. **Join:** Perform a hash join on **O_ORDERKEY** from **orders** and **L_ORDERKEY** from **lineitem**.
3. **Output:** Return the joined rows.

Question 2: What is the Impact of Selective Data Retrieval on Query Performance in Join Operations?

This question examines how fetching only necessary columns rather than all columns influences the performance of join operations.


```
-- Query fetching all columns
SELECT *
FROM orders
JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY;
```

```
-- Optimized query fetching only necessary columns
SELECT      orders.O_ORDERKEY,      lineitem.L_ORDERKEY,      orders.O_TOTALPRICE,
lineitem.L_QUANTITY, lineitem.L_SHIPDATE
FROM orders
JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY;
```

TRC Representation:

```
{ t.O_ORDERKEY, t.L_ORDERKEY, t.O_TOTALPRICE, t.L_QUANTITY, t.L_SHIPDATE | ∃ o ∈
orders, ∃ l ∈ lineitem (o.O_ORDERKEY = l.L_ORDERKEY ∧ t = o ⋈ l) }
```

Query Plan:

1. **Scan:** Scan both **orders** and **lineitem** tables.
2. **Join:** Perform a hash join on **O_ORDERKEY** from **orders** and **L_ORDERKEY** from **lineitem**.
3. **Project:** Select only **O_ORDERKEY**, **L_ORDERKEY**, **O_TOTALPRICE**, **L_QUANTITY**, and **L_SHIPDATE**.
4. **Output:** Return the selected columns.

Question 3: How Efficient are Join Operations on Filtered Data Sets Compared to Unfiltered Data Sets?

This question explores the efficiency of join operations when applied to filtered datasets as opposed to complete, unfiltered datasets.

```
-- Unfiltered join operation
SELECT *
FROM orders
JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY;
```

```
-- Join operation with a filter applied
SELECT *
FROM orders
JOIN lineitem ON orders.O_ORDERKEY = lineitem.L_ORDERKEY
WHERE lineitem.L_SHIPDATE >= '1995-04-25' AND lineitem.L_SHIPDATE < '1996-02-01';
```

TRC Representation Unfiltered Join:

```
{ t | ∃o ∈ orders, ∃l ∈ lineitem (o.O_ORDERKEY = l.L_ORDERKEY ∧ t = o ⋈ l) }
```

TRC Representation Filtered Join:

```
{ t | ∃o ∈ orders, ∃l ∈ lineitem (o.O_ORDERKEY = l.L_ORDERKEY ∧ l.L_SHIPDATE ≥ '1995-04-25' ∧ l.L_SHIPDATE < '1996-02-01' ∧ t = o ⋈ l) }
```

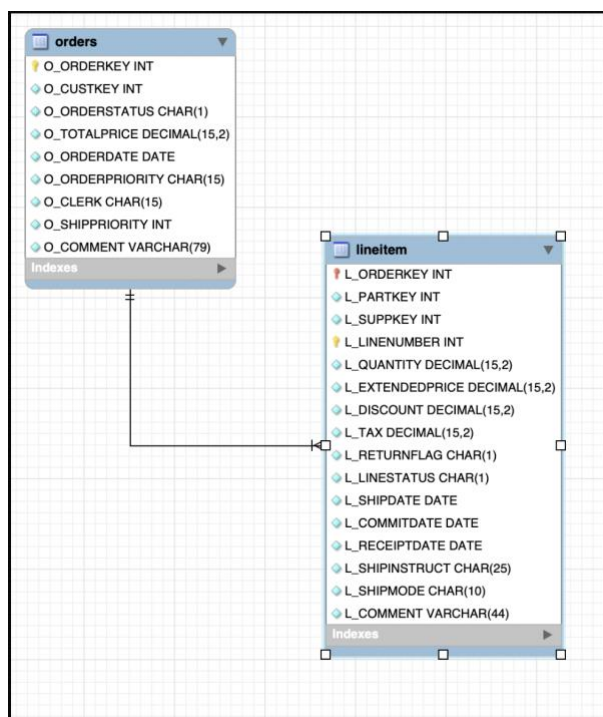
Query Plan for Unfiltered Join:

1. **Scan:** Scan both **orders** and **lineitem** tables.
2. **Join:** Perform a hash join on **O_ORDERKEY** from **orders** and **L_ORDERKEY** from **lineitem**.
3. **Output:** Return the joined rows.

Query Plan for Filtered Join:

1. **Scan:** Scan both **orders** and **lineitem** tables.
2. **Filter:** Apply filter on **L_SHIPDATE** in the **lineitem** table.
3. **Join:** Perform a hash join on **O_ORDERKEY** from **orders** and **L_ORDERKEY** from **lineitem**.
4. **Output:** Return the joined rows.

ERR Diagram:



Relationship:

The ERD shows a relationship between **orders** and **lineitem**, indicated by the connector between them. This signifies that each order can have multiple line items, establishing a one-to-many relationship where **L_ORDERKEY** in the **lineitem** table is a foreign key referencing **O_ORDERKEY** in the **orders** table.

Limitations and Constraints of the Database Performance Optimization Project

In our project, which focused on optimizing database performance with enhanced hash join algorithms using TPC-H data, we encountered several limitations and constraints. These challenges shaped the scope and outcomes of our project.

Data Limitations

1. Dataset Size and Scalability:

- Our project primarily utilized the TPC-H dataset, which, while comprehensive, has its limitations in terms of variety and real-world complexity. The data size, although significant, may not fully represent extremely large datasets found in some industrial applications.
- Scalability testing was constrained by the dataset sizes provided by TPC-H's **dbgen** tool. Extending our findings to much larger datasets involves assumptions that may not hold in all cases.

2. Data Diversity:

- The TPC-H dataset, designed for benchmarking, lacks the diversity and unpredictability of real-world data. Issues like irregular data patterns, anomalies, and real-time data variations weren't fully represented.
- The absence of such data characteristics might limit the applicability of our findings to more complex, real-world scenarios.

Project Scope Boundaries

1. Focus on Specific Hash Join Algorithms:

- The project specifically targeted optimizing hash join algorithms. While this focus allowed for in-depth analysis and optimization, it also meant that other potential areas of database performance enhancement were not explored.
- The project did not cover other types of joins (like nested loop joins or merge joins) that could be relevant in certain database contexts.

2. Use of Specific Technologies:

- Our project's optimizations were demonstrated using SQL and Python. This choice may not encompass the diverse range of database technologies and programming languages used in different environments.
- The integration of advanced hashing functions was limited to Python, which, while effective, is not a direct substitute for database-level optimization.

Challenges Encountered

1. Integration of Advanced Hashing Functions:

- A significant challenge was integrating advanced hashing functions like XXHash and FarmHash. SQL's limited support for these functions led us to use Python, introducing a layer of complexity in comparing these results directly with SQL-based operations.
- The results obtained in Python, though valuable, had to be interpreted carefully to understand their implications within SQL environments.

2. Balancing Optimization and Practicality:

- Achieving the optimal balance between performance enhancement and practical implementation was challenging. Some optimizations that theoretically promised performance gains proved to be less effective in practice due to underlying database complexities.
- The project also had to navigate the trade-offs between execution time, fetch time, and overall resource utilization.

Suggestions for Future Research and Extensions

1. Exploring Larger and More Diverse Datasets:

- Future work could involve testing with larger and more diverse datasets, perhaps even custom-generated data that includes more real-world anomalies and irregularities. This would provide a more robust test of the optimizations.

2. Extending to Other Types of Joins:

- Expanding the project to include optimizations for other types of join operations, like nested loop joins or merge joins, could provide a more comprehensive view of database performance tuning.

3. Direct Integration of Advanced Hashing in Databases:

- Investigating ways to directly integrate advanced hashing functions into database systems could be a valuable extension. This might involve working with open-source databases or developing custom database extensions.

4. **Testing in Varied Database Environments:**

- Conducting tests across different database management systems (DBMS) could reveal how optimizations perform in varied environments, offering insights that are more universally applicable.

5. **Real-Time Data and Transactional Testing:**

- Focusing on real-time data processing and transactional scenarios could add another dimension to the project, testing the optimizations in more dynamic and demanding environments.

6. **Machine Learning for Query Optimization:**

- Applying machine learning techniques to predict the most effective optimizations based on query characteristics and data patterns could be a novel area to explore.

External Sources Used in the Project

1. **TPC-H Benchmark Documentation:**

- The official TPC-H Benchmark documentation provided by the Transaction Processing Performance Council (TPC) was instrumental in understanding and implementing the benchmark tests.
- Website: https://www.tpc.org/tpc_documents_current_versions/current_specifications5.asp

2. **TPC-H Data Generation Tool (dbgen):**

- The **dbgen** tool from TPC-H was used to generate the test data sets. This tool is essential for creating a standardized and controlled environment for performance testing.
- Download Link: https://www.tpc.org/tpc_documents_current_versions/current_specifications5.asp